

Linear Regression

Steps Include

- 1. Import Libraries
- 1. Import Dataset and Preprocessing
- 1. Exploratory Data Analysis
- 1. Divide the dataset into test and train
- 1. Training the Linear Regression Model
- 1. Testing the Linear Regression Model
- 1. Evaluate Predictions

1. Import Libraries

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

2. Import Dataset and Preprocessing

```
In [2]: df=pd.read_csv("IceCreamData.csv")
df.head()
```

```
Out[2]:
```

	Temperature	Revenue
0	24.566884	534.799028
1	26.005191	625.190122
2	27.790554	660.632289
3	20.595335	487.706960
4	11.503498	316.240194

```
In [3]: df.shape
```

```
Out[3]: (500, 2)
```

```
In [4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 500 entries, 0 to 499
Data columns (total 2 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   Temperature     500 non-null   float64
 1   Revenue         500 non-null   float64
dtypes: float64(2)
memory usage: 7.9 KB
```

```
In [12]: df.isnull().sum()
```

```
Out[12]: Temperature    0  
Revenue              0  
dtype: int64
```

```
In [9]: df.describe().round(2)
```

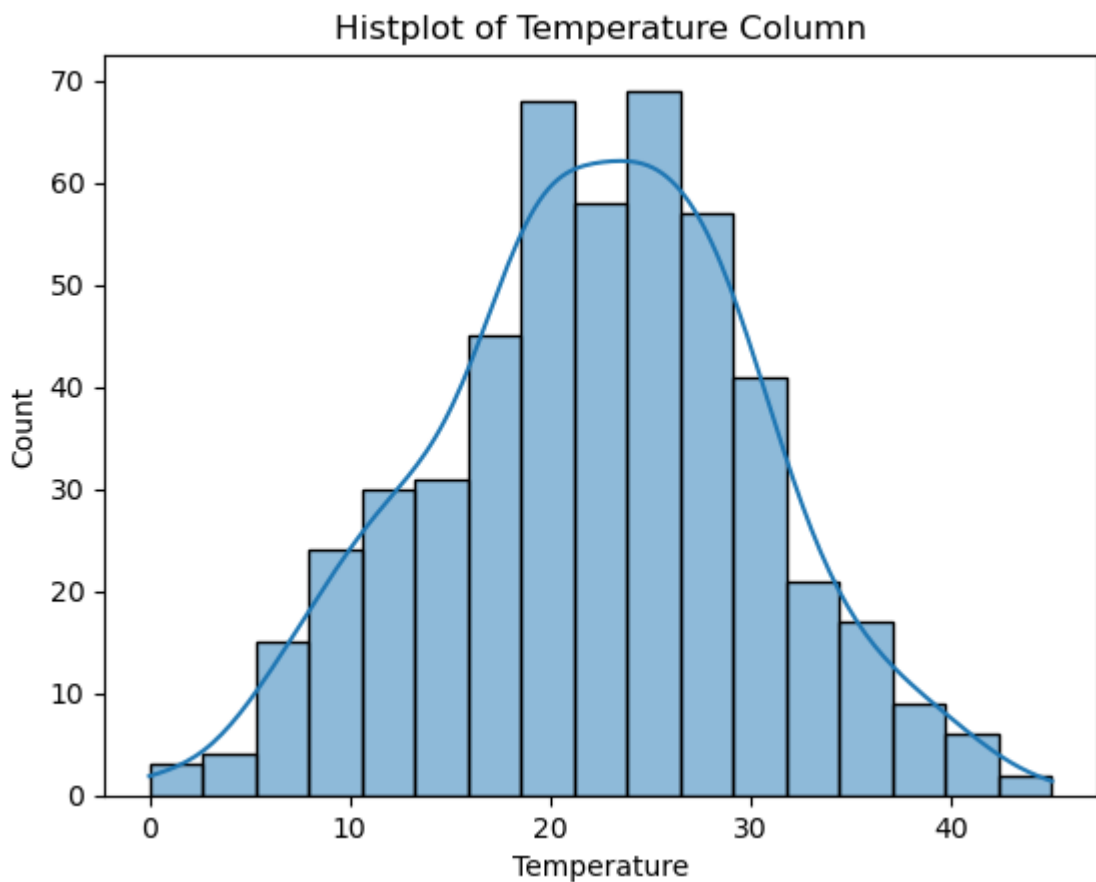
```
Out[9]:
```

	Temperature	Revenue
count	500.00	500.00
mean	22.23	521.57
std	8.10	175.40
min	0.00	10.00
25%	17.12	405.56
50%	22.39	529.37
75%	27.74	642.26
max	45.00	1000.00

3. Exploratory Data Analysis

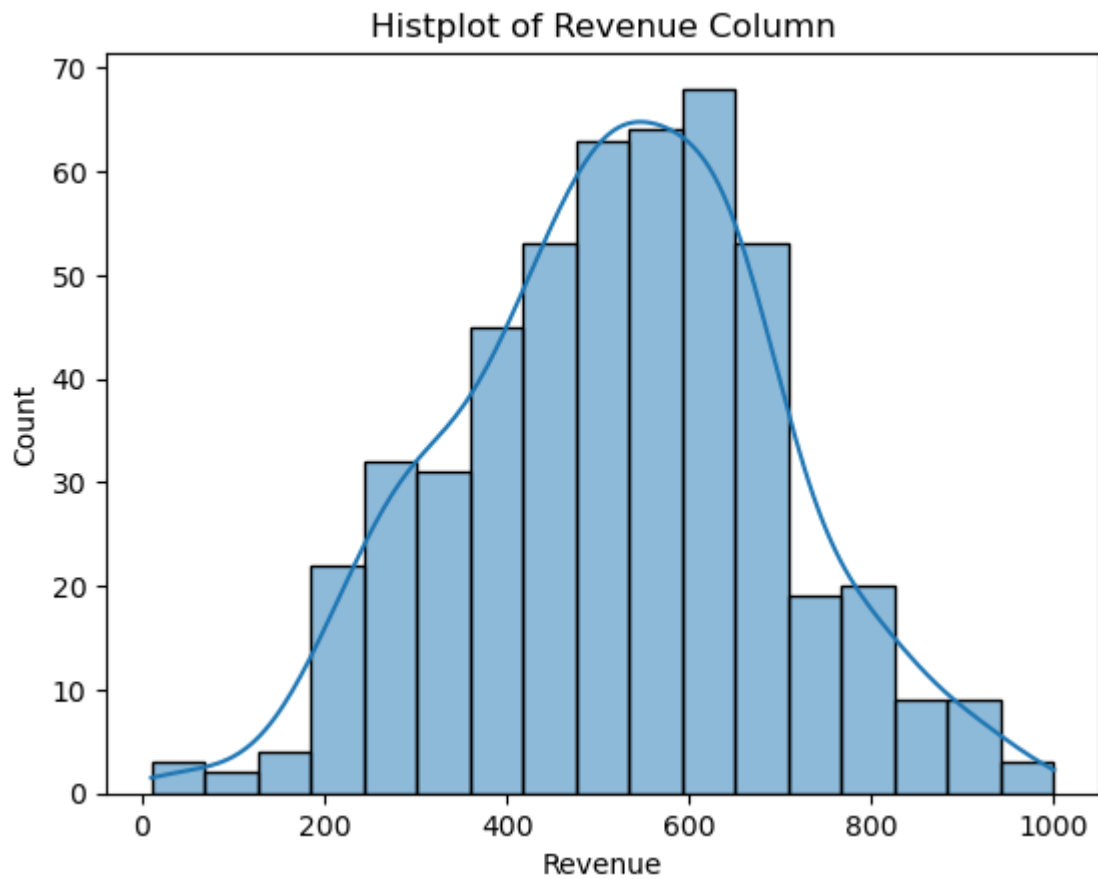
```
In [83]: plt.title("Histogram of Temperature Column")  
sns.histplot(data=df, x="Temperature", kde=True)
```

```
Out[83]: <Axes: title={'center': 'Histogram of Temperature Column'}, xlabel='Temperature', ylabel='Count'>
```




```
In [84]: plt.title("Histogram of Revenue Column")  
sns.histplot(data=df, x="Revenue", kde=True)
```

Out[84]: <Axes: title={'center': 'Histplot of Revenue Column'}, xlabel='Revenue', ylabel='Count'>




```
In [87]: # plt.title("Scatter Plot between Temperature and Revenue")
sns.relplot(data=df, x="Temperature", y="Revenue")
```

Out[87]: <seaborn.axisgrid.FacetGrid at 0x16932b090>




```
In [11]: sns.relplot(data=df,x="Temperature",y="Revenue",kind="line")
```

```
Out[11]: <seaborn.axisgrid.FacetGrid at 0x155db8710>
```



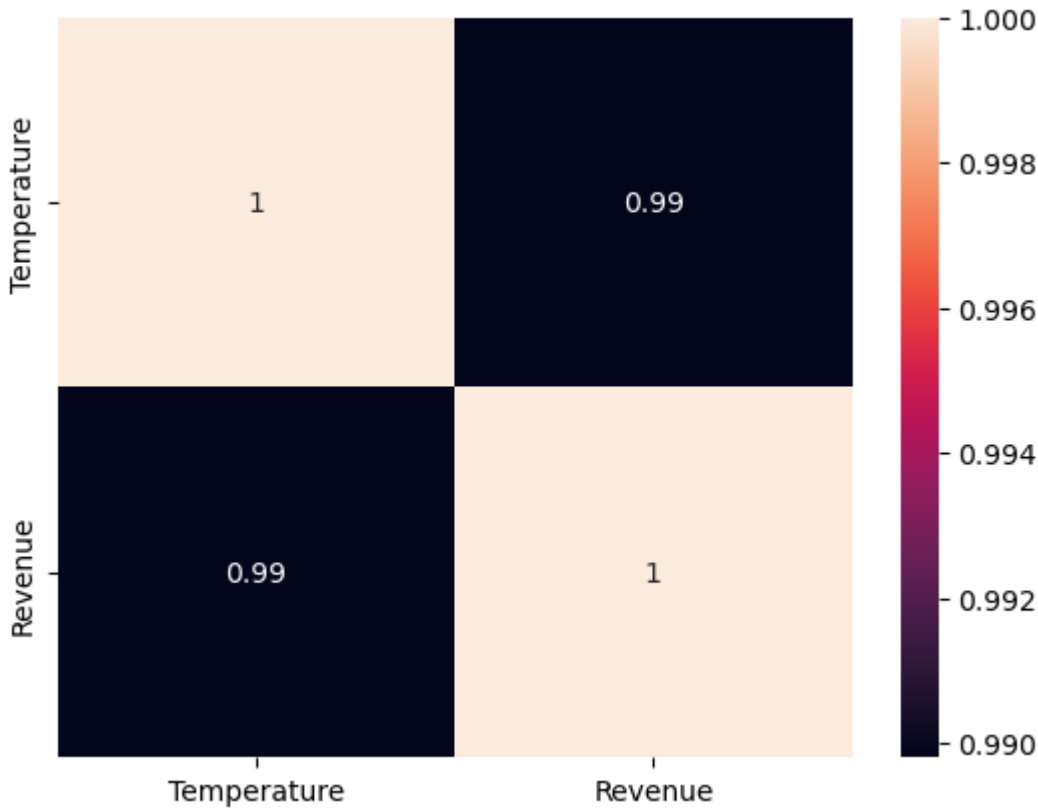
```
In [91]: sns.pairplot(df)
```

```
Out[91]: <seaborn.axisgrid.PairGrid at 0x16b1ef350>
```



```
In [15]: corr=df.corr()  
sns.heatmap(corr,annot=True)
```

Out[15]: <Axes: >



4. Divide the Dataset into Train and Test

```
In [16]: from sklearn.model_selection import train_test_split
```

```
In [17]: df.head()
```

```
Out[17]:
```

	Temperature	Revenue
0	24.566884	534.799028
1	26.005191	625.190122
2	27.790554	660.632289
3	20.595335	487.706960
4	11.503498	316.240194

```
In [49]: df_numpy=np.array(df)
```

```
In [50]: df_numpy.shape
```

```
Out[50]: (500, 2)
```

```
In [52]: x=df_numpy[:, :-1]
         y=df_numpy[:, -1]
```

```
In [53]: x.shape,y.shape
```

```
Out[53]: ((500, 1), (500,))
```

```
In [55]: x_train.shape
```

```
Out[55]: (400, 1)
```

```
In [56]: x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.2)
```

```
In [57]: x_train.shape,x_test.shape,y_train.shape,y_test.shape
```

```
Out[57]: ((400, 1), (100, 1), (400,), (100,))
```

5. Training the Linear Regression Model

```
In [59]: from sklearn.linear_model import LinearRegression
```

```
In [61]: le=LinearRegression()
```

```
In [62]: le.fit(x_train,y_train)
```

```
Out[62]:
```

▼ LinearRegression
 LinearRegression()

6. Testing the Linear Regression Model


```
In [63]: y_pred=le.predict(x_test)
```

```
In [69]: le.score(x_test,y_test)
```

```
Out[69]: 0.9809948176474016
```

```
In [82]: plt.scatter(x_train,y_train)
plt.plot(x_test,y_pred,color="red")
```

```
Out[82]: [matplotlib.lines.Line2D at 0x16a8539d0>]
```



7. Evaluate Predictions

```
In [66]: from sklearn.metrics import accuracy_score,mean_squared_error,mean_absolute_
```

```
In [76]: Acc = le.score(x_test,y_test)
```

```
In [72]: MAE = mean_absolute_error(y_test,y_pred)
```

```
In [73]: MSE = mean_squared_error(y_test,y_pred)
```

```
In [74]: RMSE = np.sqrt(MSE)
```

```
In [77]: R2 = r2_score(y_test,y_pred)
```

```
In [92]: print("Accuracy Score is : ",Acc)
print("Mean Absolute Error is : ",MAE)
print("Mean Squared Error is : ",MSE)
```



```
print("Root Mean Squared Error is : ",RMSE)  
print("R2 Score is : ",R2)
```

```
Accuracy Score is : 0.9809948176474016  
Mean Absolute Error is : 19.47569835356527  
Mean Squared Error is : 576.6401232652831  
Root Mean Squared Error is : 24.013332198286918  
R2 Score is : 0.9809948176474016
```